

CSE 447: Assignment 1

Abosh Upadhyaya *

January 26, 2024

Question 1: n -gram Language Models

Part 1: n -gram Language Modeling & Perplexity

(a) Describe how you built your n -gram models. Provide graphs, tables, charts or other summary evidence to support any claims you make.

I built my n -gram models using a class called `NGramModel` which has a couple fields and methods. The fields are:

- `n`: The n in n -gram.
- `ngram_counts`: A dictionary that maps n -grams to their counts.
- `context_counts`: A dictionary that maps contexts to their counts.
- `vocabulary`: A set of all the words in the vocabulary.
- `smoothing_constant`: The smoothing constant used for Laplace smoothing. Defaults to 1.
- `lambdas`: A list of the λ values used for interpolation (for the trigram model).

From here, implementing the methods wasn't that bad. I just went back to the slides and implemented the necessary methods to complete the training and prediction processes.

To train the model, I first preprocessed the data. This entailed tokenizing it (naive whitespace tokenization) and adding $n - 1$ START tokens to each line and a STOP token to the end of each line. I also converted all the tokens that appeared less than 3 times to UNK tokens.

Then, I built `vocabulary` by iterating through each line and adding each token to the set, making sure to discard the START token.

Finally, I updated the `ngram_counts` and `context_counts` dictionaries by iterating through each n -gram in each line, and updating the counts accordingly.

To predict (and calculate the model perplexity), I first preprocessed the data that needed to be predicted. This entailed tokenizing it (the same procedure as before, except this time we only UNKify tokens that appear less than 3 times in the *training* data).

Then, for each line, I iterated through each n -gram in the line and used the maximum likelihood estimate to calculate the probability of the n -gram given the context. I then took the log of this probability and added it to a running sum of the log probabilities of each n -gram in the line.

***Collaborators: ChatGPT.** ChatGPT helped by providing debugging advice, as well as explaining how to use inheritance in Python (for the `WordPieceTokenizer` which inherited from the `BytePairTokenizer`). It also helped me with making the scatterplots (using `matplotlib`) and formatting the L^AT_EX in this document.

If the user wanted to use Laplace smoothing (specified by the `use_laplace_smoothing` parameter), I simply added the `smoothing_constant` to the numerator and the `smoothing_constant` times the vocabulary size to the denominator.

Finally, I returned the perplexity of the model by taking the exponent of the negative of the average log probability of each n -gram in the line. The average was calculated by dividing the total summed log probabilities for all sentences in the set divided by the total number of tokens in the set, making sure to subtract the amount of START tokens in each line, as they're not part of the vocabulary and our model can't predict them.

Adding interpolation was not that bad. I simply added a new method called `perplexity_with_interpolation` which did almost exactly the same thing, except it also took a unigram and bigram model (to be used for interpolation). From there, I computed the average perplexity by doing the same thing as before except using the interpolation formula.

(b) Describe how you computed *perplexity* without any smoothing in a detailed equation, in natural language, or with pseudo code.

As I stated in my answer to part (a), I returned the perplexity of the model by taking the exponent of the negative of the average log probability of each n -gram in the line. The average was calculated by dividing the total summed log probabilities for all sentences in the set divided by the total number of tokens in the set, making sure to subtract the amount of START tokens in each line, as they're not part of the vocabulary and our model can't predict them.

(c) What's the potential downside of Laplace smoothing? Back up your claim with empirical evidence.

One potential downside of Laplace smoothing is that it might add too much probability mass to a rare event or n -gram. Which is kind of crazy, because it's only adding 1 to the count of each token. But, if you think about it, if you have a very large vocabulary, adding 1 to the count of each token will add a lot of probability mass to the rare or unseen tokens, which barely appear at all in the training data.

Basically, because it's uniform, it adds too much probability mass to the rare events, which is not good.

Here's some empirical evidence to show why this is bad. I analyzed the unigram, bigram, and trigram models with and without Laplace smoothing on **the training set**. Here were the results:

Model	No Smoothing	Laplace Smoothing ($k = 1.0$)
1-gram	976.54	977.50
2-gram	77.07	1,442.30
3-gram	7.87	6,244.42

Table 1. Perplexity scores of language models with and without Laplace smoothing ($k = 1.0$)

As can be seen, the perplexity of the models with Laplace smoothing is significantly higher than the perplexity of the models without Laplace smoothing **on the training set**.

(d) Another extension of Laplace smoothing, instead of adding 1 to the count of each token, is to add k , where typically $0 < k < 1$. Try different values for k , and describe how different k change the perplexity score differently.

I ran the same experiment, except this time with a smoothing constant of $k = 0.001$. Here were the results:

Model	No Smoothing	Laplace Smoothing ($k = 0.001$)
1-gram	976.54	976.54
2-gram	77.07	95.08
3-gram	7.87	35.09

Table 2. Perplexity scores of language models with and without Laplace smoothing ($k = 0.001$)

Of course, no smoothing is best on the training set as expected since our model was trained on the training set. But, as k got smaller and smaller, the perplexity got smaller and smaller on the **training set**. This makes sense, because as k gets smaller and smaller, the amount of probability mass added to each token gets smaller and smaller, which is better for the model.

However, on the dev and test sets, the happy medium was not just to move k closer and closer to 0. In fact, the best k for the dev set was somewhere around 0.03 (although I didn't rigorously test this, I just plugged some values in and saw the best one, so maybe there's a better one.).

(e) Report the *unsmoothed* and *smoothed with Laplace smoothing* version of perplexity scores of the unigram, bigram, and trigram language models for your training, development, and test sets. Briefly discuss the experimental results.

Here are the results:

Model	No Smoothing	Laplace Smoothing ($k = 1.0$)
1-gram (Training)	976.54	977.51
2-gram (Training)	77.07	1,442.30
3-gram (Training)	7.87	6,244.42
1-gram (Dev)	892.24	894.39
2-gram (Dev)	∞	1,669.65
3-gram (Dev)	∞	9,676.65
1-gram (Test)	896.49	898.55
2-gram (Test)	∞	1,665.38
3-gram (Test)	∞	9,649.60

Table 3. Perplexity scores of language models with and without Laplace smoothing ($k = 1.0$)

It is expected that the model had infinite perplexity on the dev and test sets. This is because the model was trained on the training set, and the dev and test sets might contain n-grams that were never seen in the training sets. Thus, the model would assign a probability of 0 to these n-grams, and the perplexity would be infinite.

It's also interesting to see the perplexity be so high for $k = 1.0$. As I said in my answer to part (d),

if we change k to be something like 0.03, these are the results:

Model	No Smoothing	Laplace Smoothing ($k = 0.03$)
1-gram (Training)	976.54	976.54
2-gram (Training)	77.07	235.46
3-gram (Training)	7.87	409.10
1-gram (Dev)	892.25	892.29
2-gram (Dev)	∞	521.28
3-gram (Dev)	∞	3,567.91
1-gram (Test)	896.50	896.54
2-gram (Test)	∞	519.04
3-gram (Test)	∞	3,549.07

Table 4. Perplexity scores of language models with and without Laplace smoothing ($k = 0.03$)

Notice how now, the perplexity for the bigram model is much lower than the perplexity for the unigram model! This could be because of the data of course, but it was an interesting result.

(f) Your goal is to find reasonably good combinations of $\lambda_1, \lambda_2, \lambda_3$. Experiment and report perplexity scores on training and development sets for five sets of values of $\lambda_1, \lambda_2, \lambda_3$ that you tried, along with short descriptions of the strategies that you used to find better hyperparameters. In addition, report the training and development perplexity for the values $\lambda_1 = 0.1, \lambda_2 = 0.3, \lambda_3 = 0.6$.

Here are the results of running grid search to find the best λ values for interpolation on the training and dev sets:

λ values	Training Perplexity	Dev Perplexity
(0.1, 0.1, 0.8)	9.33	482.85
(0.1, 0.3, 0.6)	11.15	352.23
(0.1, 0.5, 0.4)	14.17	308.64
(0.3, 0.1, 0.6)	11.97	364.09
(0.3, 0.3, 0.4)	15.30	286.63
(0.3, 0.5, 0.2)	22.44	262.89
(0.5, 0.1, 0.4)	16.82	334.72
(0.5, 0.3, 0.2)	25.06	280.50

Table 5. Perplexity scores for different λ values on an interpolated trigram model

I determined the optimal hyperparameters with a standard grid search. I tried a bunch of different combinations of λ values that summed to 1, and printed their perplexity scores on the training and dev sets. I then picked the λ values that had the lowest perplexity on the dev set.

(g) Putting it all together, report perplexity on the test set, using the best combination of hyperparameters that you chose from the development set. Specify those hyperparameters.

λ values	Test Perplexity
(0.3, 0.5, 0.2)	262.66

Table 6. Perplexity for the best λ values on an interpolated trigram model

(h) If you use half of the training data, would it increase or decrease the perplexity of previously unseen data? Why? Provide empirical experimental evidence to support your claims.

It will increase the perplexity of previously unseen data. This is because the model will be trained on less data, so it will be less accurate. This will cause the perplexity to increase, as the model has less of an "idea" of what words could come next. This means it has to store more information each time it predicts a word, which means the perplexity will increase.

When I trained the interpolated model on half of the training data, here were the results:

λ values	Dev Perplexity	Test Perplexity
(0.1, 0.1, 0.8)	∞	∞
(0.1, 0.3, 0.6)	∞	∞
(0.1, 0.5, 0.4)	∞	∞
(0.3, 0.1, 0.6)	∞	∞
(0.3, 0.3, 0.4)	∞	∞
(0.3, 0.5, 0.2)	∞	∞
(0.5, 0.1, 0.4)	∞	∞
(0.5, 0.3, 0.2)	∞	∞

Table 7. Perplexity scores for different λ values on the dev and test sets

The perplexities on both the dev and test sets became ∞ . The only way this can happen is if the model had an interpolated probability of 0 for some n -gram in the dev or test set. That means it encountered some unseen token in the dev or test set, which is more likely now that we only trained the model on half of the training data.

Normally (in this assignment), at least the unigram model will have a non-zero probability for an n -gram since all individual tokens in the dev and test set are in the training set, but because there's no smoothing with the interpolated model, a single unseen token will cause the interpolated probability to be 0, which will cause the perplexity to be ∞ . And since we only trained on half the training data, there were some unseen tokens.

(i) If you convert all tokens that appeared less than five times to $\langle \text{unk} \rangle$ (a special symbol for out-of-vocabulary tokens), would it increase or decrease the perplexity on the previously unseen data compared to an approach that converts only those words that appeared just once to $\langle \text{unk} \rangle$? Why? Provide empirical evidence to support your claims.

Converting all tokens that appear less than 5 times to $\langle \text{UNK} \rangle$ will decrease the perplexity on the previously unseen data compared to an approach that converts only those words that appeared just 1 to $\langle \text{UNK} \rangle$.

Here are the results after testing different λ values on the dev and test sets with different $\langle \text{UNK} \rangle$ thresholds:

λ values	Dev (= 1)	Dev (< 5)	Test (= 1)	Test (< 5)
(0.1, 0.1, 0.8)	559.92	391.15	555.84	389.67
(0.1, 0.3, 0.6)	406.70	287.70	403.94	286.85
(0.1, 0.5, 0.4)	355.51	253.19	353.33	252.64
(0.3, 0.1, 0.6)	415.13	302.13	413.20	301.48
(0.3, 0.3, 0.4)	325.14	239.99	323.75	239.63
(0.3, 0.5, 0.2)	297.28	221.23	296.22	221.11
(0.5, 0.1, 0.4)	378.09	281.50	377.00	281.24
(0.5, 0.3, 0.2)	315.07	238.08	314.33	238.09

Table 8. Perplexity scores for different λ values with different $\langle \text{UNK} \rangle$ thresholds

This decrease in perplexity is probably because the model is dealing with a smaller, more generalized vocabulary. In other words, the model is better at handling unseen data when it generalizes words that don't appear frequently into the $\langle \text{UNK} \rangle$ token, rather than treating them as individual tokens. This generalization probably allows the model to use its learned probabilities more effectively because it isn't distracted by the noise of rare words, which it hasn't learned meaningful information about during training.

Question 2: Byte-Pair Encoding

(a) Please produce a scatterplot showing points (x, y) , each corresponding to an iteration of the algorithm, with x the current size of the type vocabulary (including the base vocabulary), and y the length of the training corpus (in tokens) under that vocabulary's types. How many types do you end up with? What is the length of the training data under your final type vocabulary?




Figure 1. Scatterplot of training corpus length (in tokens) (y -axis) versus vocabulary size (x -axis). Each point represents the state at every 500th iteration of the BPE training algorithm.

The final vocabulary size was 9,219 types. The number of tokens in the training corpus after termination was 101,849.

(b) Another way of overcoming the rare word issue is to encode text as a sequence of characters. Discuss the advantages and potential issues of character-level encoding compared to BPE.

Advantages over BPE	Disadvantages to BPE
This is a much simpler algorithm to implement. It's just a simple loop that iterates through each character in the training data and adds it to the vocabulary. This is much simpler than the BPE algorithm, which can be really complicated to implement and debug.	Because we're encoding each character, the length of a sequence the model has to process is much longer . This leads to worse performance and efficiency. BPE however, can encode a sequence of characters as a single token, which is much more efficient.
There is no need to worry about out-of-vocabulary words , since we are essentially encoding every single possible character in the language. Thus, our model can represent any text, regardless of whether it's in the training data or not. BPE on the other hand, can only represent tokens that are in the training data, so there's still a possibility of encountering an out-of-vocabulary token.	Linguistically speaking, there is a benefit in storing tokens that aren't just characters. This is because words have meaning , and encoding them as individual characters can lose that meaning. For example, the suffix "ing" is a common suffix in English that denotes the present participle tense. If we encode each character as a token, we lose this meaning, and the model has to learn it from scratch. BPE allows us to retain this meaning, as it can encode the suffix "ing" as a single token.

Table 9. Advantages and disadvantages of character-level encoding compared to BPE.

(c) Applying your tokenizer on the last 1000 lines of the data, how many tokens do you end up with? Do you encounter issues with words that didn't appear in the training set? More generally, when would you expect your tokenizer to fail when applying to an unseen text.

I ended up with 27,969 tokens. I did not encounter any issues with words that didn't appear in the training set. I'm guessing this is because the training and test sets are from the same language distribution, so they are correlated well and the tokens BPE formed matched the test set well.

However, in general, this is true of BPE anyways. BPE picks up on linguistic patterns inside of words themselves and forms tokens based on those patterns. But, there are obviously cases where BPE can fail.

An easy example is if the training data is not representative of the language patterns in the test set or in the real world (depending on your model's use case). For example, if the training data is from a different language distribution than the test set, then BPE will fail because it will form tokens based on the training data, which will not match the test set. It'd be like teaching someone to speak in a certain dialect which is region-specific, and then asking them to try and predict what someone who was from a completely different part of the world would say. They'd have no knowledge of the language patterns in that region, so they'd do a bad job even if they understood their dialect well.

Question 3: WordPiece

(a) Please produce a scatterplot showing points (x, y) , each corresponding to an iteration of the algorithm, with x the current size of the type vocabulary (including the base vocabulary), and y the length of the training corpus (in tokens) under that vocabulary's types. How many types do you end up with? What is the length of the training data under your final type vocabulary?




Figure 2. Scatterplot of training corpus length (in tokens) (y-axis) versus vocabulary size (x-axis). Each point represents the state at every 100th iteration of the WordPiece training algorithm.

The final vocabulary size was 4,083 types. The number of tokens in the training corpus after termination was 455,238 (hey, less tokens means we merged more!).

(b) Applying your tokenizer on the last 1000 lines of the data, report the length of the tokenized data. Also, include the tokenized sequences for the following two sentences:

The length of the tokenized data in tokens was 108,692 tokens.

i) "Analysts were expecting the opposite, a deepening of the deficit."

['A', 'n', 'a', 'l', 'y', 's', 't', 's', 'G',
'w', 'e', 'r', 'e', 'G',

‘exp’, ‘e’, ‘c’, ‘t’, ‘i’, ‘n’, ‘g’, ‘Ġ’,
 ‘th’, ‘e’, ‘Ġ’,
 ‘o’, ‘p’, ‘p’, ‘o’, ‘s’, ‘i’, ‘t’, ‘e’, ‘,’’, ‘Ġ’,
 ‘a’, ‘Ġ’,
 ‘d’, ‘e’, ‘e’, ‘p’, ‘e’, ‘n’, ‘i’, ‘n’, ‘g’, ‘Ġ’,
 ‘o’, ‘f’, ‘Ġ’,
 ‘d’, ‘e’, ‘f’, ‘i’, ‘c’, ‘i’, ‘t’, ‘.’]

ii) *“Five minutes later, a second person arrived, aged around thirty, with knife wounds.”*

[‘F’, ‘i’, ‘v’, ‘e’, ‘Ġ’,
 ‘m’, ‘i’, ‘n’, ‘u’, ‘t’, ‘e’, ‘s’, ‘Ġ’,
 ‘l’, ‘a’, ‘t’, ‘e’, ‘r’, ‘,’’, ‘Ġ’,
 ‘a’, ‘Ġ’,
 ‘s’, ‘e’, ‘c’, ‘o’, ‘n’, ‘d’, ‘Ġ’,
 ‘p’, ‘e’, ‘r’, ‘s’, ‘o’, ‘n’, ‘Ġ’,
 ‘a’, ‘r’, ‘r’, ‘i’, ‘v’, ‘e’, ‘d’, ‘,’’, ‘Ġ’,
 ‘a’, ‘g’, ‘e’, ‘d’, ‘Ġ’,
 ‘a’, ‘r’, ‘o’, ‘u’, ‘n’, ‘d’, ‘Ġ’,
 ‘th’, ‘i’, ‘r’, ‘t’, ‘y’, ‘,’’, ‘Ġ’,
 ‘w’, ‘i’, ‘th’, ‘Ġ’,
 ‘k’, ‘n’, ‘i’, ‘f’, ‘e’, ‘Ġ’,
 ‘w’, ‘o’, ‘u’, ‘n’, ‘d’, ‘s’, ‘.’, ‘.’]

(c) In terms of efficiency and performance, what’s the advantages and disadvantages of WordPiece compared with BPE? There is no single correct answer here, just provide your thoughts and rationales, supported by empirical evidences.

Advantages over BPE	Disadvantages to BPE
WordPiece normalizes the frequency of the two adjacent tokens by their relative frequencies. For example, if the tokens "Hi there" appear 100 times together in the training set, much higher than any other adjacent tokens, BPE would immediately merge them together as a token and move forward. Meanwhile, if the tokens "Calif" and "ornia" appear only 5 times together in the training set, BPE would skip them on this iteration (and likely) continue to skip them for many iterations, since they’re not frequently occurring together. However, WordPiece would note that yes, they don’t appear together often, but they almost only appear <i>together</i> . That is, you never see "Calif" or "ornia" by itself, and you do see them together, which implies they should be merged. So WordPiece picks up on these linguistic patterns and merges them together, which is a huge advantage.	BPE is simpler to understand than WordPiece. The tokenization algorithm for WordPiece can also be difficult, since we’re searching for the longest matching subword, instead of just merging tokens until we can’t merge them anymore! Additionally, the optimization criteria for BPE is simple: simply choose the most frequently occurring pair of adjacent tokens and merge them. WordPiece is a bit more complex, and furthermore, requires more calculations as we have to not only calculate the frequency of the pair of tokens but also the frequency of each token individually. This makes it more computationally expensive than BPE.

Table 10. Advantages and disadvantages of WordPiece compared to BPE.

In terms of **empirical evidence**, if we look at the two graphs, it isn’t directly clear which one is better. Sure, they both decrease the size of the training corpus as they merge tokens together, but we can’t just compare the size of the test corpus to determine which one is better, because we had different training criterion.

Perhaps the log scale of the BPE graph is because we trained for too long, and after a while, it becomes harder and harder to minimize the size of the training corpus. And perhaps we trained the WordPiece model for too short of a time, and if we trained it for longer, we’d see that sharp decrease at the end for even longer at the end of the graph.

In BPE, we stopped when the frequency of the next most frequent pair was 2. In WordPiece, we just learned 4,000 merge rules and stopped. So it’s not fair to compare these models as is.

However, I do know about WordPiece’s usage with BERT, and I know that it performs well with masked-language modeling (where you mask out a word and try to predict it). BPE might excel at tasks that require more common phrases, as it can more efficiently encode them, but I think that WordPiece would perform similarly as well, as it’s also still directly correlated to the frequency of the tokens appearing together! In a sense, WordPiece is just a level up from BPE, as it’s more sophisticated and requires more calculations, but it’s still based on the same idea of merging tokens together based on their frequency of occurrence.