

CSE 447: Assignment 1

Abosh Upadhyaya *

February 26, 2024

Question 1: Baseline Evaluation

(a) Run the pretrained model on the evaluation function over the entire test set. What is the average reward?

The average reward is `0.6231`.

(b) Please provide 5 example generations from the model. How do you feel about the quality of the generations? What sentiment do they express? Feel free to say so in your report if the sentences don't express clear sentiment to you.

Five example generations are:

1. "Young Elijah Wood and Joseph Mazzello are outstanding musicians and performers for the United States and the world.\n\n\n\n\nFollow @J"
2. "This is one of the best sequels around and a real star, just because they are better than many others.\n\n\n\nYou may not know"
3. "That's what me and my friends kept asking each other the day about. At one point, I asked if she'd take some of it away from"
4. "'Mame' is a disgrace to many things he did.\n\n P.S. I don't actually like it too much...\nM"
5. "The Tender Hook, or, Who Killed The Tender Hook, or, Who Killed The Tender Hook, or, Who Killed the Tender"

Sentences 1 and 2 express positive sentiment, even if they don't make much sense. It feels like the model doesn't really understand what the words *mean*, but can pick up on the positivity of the original tokens in the review.

Sentences 3 and 5 don't express clear sentiment at all, as they just seem like fragments of sentences that are cobbled together. Sentence 5 is just repetition, which could mean that the model is not understanding the context of the review at all.

Sentence 4 expresses negative sentiment, assuming "Mame" is a movie or something that has been reviewed.

Also, I'm overall confused by the generation of the newline characters. It seems like sometimes the model predicts many of them in a row.

***Collaborators:** ChatGPT. ChatGPT helped by providing debugging advice, as well as with formatting the L^AT_EX in this document.

(c) What does the `loop` variable do in the `evaluate` function?

The `loop` variable is an instance of the `tqdm` class, which basically displays progress bars for the evaluation process. During each iteration of the loop, the `loop` variable is updated to reflect the progress of the evaluation process. This is done by calling the `update` method of the `loop` variable, which increments the progress bar by 1. We also set its description to show the average reward calculated so far.

Overall, `loop` is used for user feedback, so we can see how far along the evaluation process is, and how the average reward is changing. The progress bar can be useful when initially testing our implementation. We can first run our cell on CPU for a few iterations, check out the progress bar to make sure everything is going as expected, and then we can re-run the cell using GPU if we're confident in our implementation. This way, we conserve GPU time and resources.

Question 2: REINFORCE Implementation

(a) Run your training function over the entire training set for 1 epoch. An efficient implementation with the recommended hyperparameters above should train the model in under 10 minutes. Plot the reward over time and report the average reward on the test set of the final trained model. A valid solution should achieve a final average reward of at least 0.75.

The average reward on the test set of the final trained model is 0.7939.




Figure 1. Reward over time.

(b) Provide 5 example generations from the model. Are they more positive than the original generations? What else do you notice about them?

Five example generations are:

1. "Young Elijah Wood and Joseph Mazzello are outstanding actors. I like this. I like this. I like this. I like this. I like"
2. "This is one of the best sequels around and a real classic,,,,,,,,,,,,,,,,,,,,,"

3. "That's what me and my friends kept asking each other the most most most most most most most most most most most popular of most most most most most"
4. "Mame" is a disgrace to many things. I like this. I like this. I like this. I like this. I like this"
5. "The Tender Hook, or, Who Killed The New York. I like this. I like this. I like this. I like this. I like this. I"

These generations are certainly more positive than the original generations, as they all contain the phrase "I like this" or some other positive phrase repeated over and over again. This is likely because the model has learned that repeating positive phrases is likely to receive a high reward.

(c) What are the limitations of REINFORCE based on your observations of the generated reviews? Explain what might have caused such limitations.

The limitations of REINFORCE based on the observations of the generated reviews are that the model seems to be stuck in a loop of repeating the same words over and over again. This is probably because the model is not learning to generate coherent sentences, and is instead trying to generate sequences of words that maximize the reward.

This is a limitation of REINFORCE because it doesn't take into account the structure of the generated sequences, and instead just tries to maximize the reward by generating sequences that are likely to receive high rewards. In this case, saying something like "I like this" over and over again is likely to receive a high reward, so it is a valid strategy for the model to pursue.

(d) What does `reset_model_optimizer` do?

The `reset_model_optimizer` function initializes / resets the pre-trained causal LM and its optimizer for training. It sets the model's padding token ID to match the end-of-sentence token ID of the tokenizer, so that the sequences are processed correctly. Also, it initializes the optimizer to an AdamW optimizer with the given learning rate.

It's basically just used to reset the model and optimizer to their initial states, so that we can train the model from scratch. This is useful when we want to train the model for multiple epochs, as we can reset the model and optimizer to their initial states after each epoch, so that we can train the model from scratch each time.

(e) What's the shape of `log_probs` in `compute_reinforce_loss`, and what does each dimension correspond to?

The shape of `log_probs` is `[batch_size, sequence_length, vocab_size]`.

In this case, the `batch_size` is 32, the `sequence_length` is 30, and the `vocab_size` is 50257.

The `batch_size` dimension corresponds to the number of sequences in the batch, the `sequence_length` dimension corresponds to the length of each sequence, and the `vocab_size` dimension corresponds to the size of the vocabulary.

Question 3: Regularization

(a) First, run your REINFORCE training function with $\alpha = 0$ for 1 epoch. This is equivalent to training without the KL-penalty. Plot the KL-divergence between the original model and the new model over time. What trend do you see from the plot? Why could such a trend be undesirable in practice?

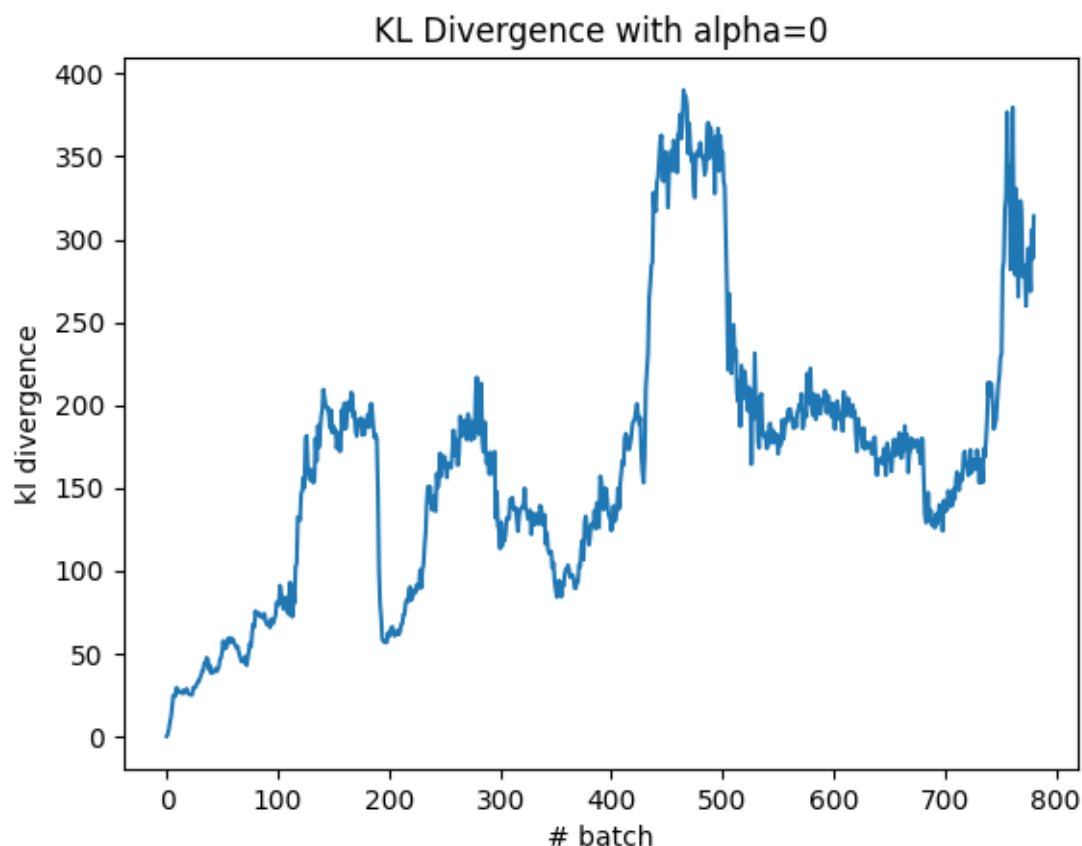


Figure 2. KL-divergence over time with $\alpha = 0$.

The trend I see from the plot is that the KL-divergence between the original model and the new model is increasing over time. This is undesirable in practice because it means that the new model is diverging from the original model, and is no longer a good approximation of the original model.

(b) Try different α to experiment with the power of KL-regularization. For each setup, report the α values, the KL-divergence plots, reward plots, 5 sample generations, and comments on the samples' quality.

A high α that prevents the model from being able to get a reward of 0.65 or more.

I tried $\alpha = 1000$.




Figure 3. KL-divergence and reward over time with $\alpha = 1000$.


Five example generations are:

1. "Young Elijah Wood and Joseph Mazzello are outstanding musicians and performers for the United States and the rest of the world. This was created in 1967 by"
2. "This is one of the best sequels around and a real star, just because they are better than many already.\n\n\nYou may see the"
3. "That's what me and my friends kept asking each other the day about. At one point, I asked if she'd take some of it away from"
4. "'Mame' is a disgrace to many things he did.\nP.S. I don't see the idea that this is a post"
5. "The Tender Hook, or, Who Killed The Tender Hook, or, Who Killed The Tender Hook, or, Who Killed the Tender"

The quality of the samples honestly looks like that in part 1. The α value is way too high, and the model seems to overly regularize which leads it to not learn anything at all. The generations don't seem positive, they just seem like fragmented natural language follow-ups, which is probably because the regularization is so high. Because of this, the reward isn't as high either, just like in part 1 where the reward was also under 0.65.

A low α such that the outputs are poor quality. I tried $\alpha = 1e8$.

Five example generations are:






Figure 5. KL-divergence and reward over time with $\alpha = 0.0003$.

Five example generations are:

1. "Young Elijah Wood and Joseph Mazzello are outstanding actors and I am a great time to do a great things. I think it is a great time"
2. "This is one of the best sequels around and a real movie, it's a great time for us and I am a great time to feel this to"
3. "That's what me and my friends kept asking each other the time. I think it's wonderful. I'm a great young, and I'm happy"
4. "'Mame' is a disgrace to many things this movie is great and we are great with no regrets at all the time. If you are a"
5. "The Tender Hook, or, Who Killed The Tenders, and I think the best time to do a great things. But I'll be very"

The quality of these samples is better! It's not perfect, and doesn't really make sense, but at least it's natural language (even if it's not grammatically correct). Things like "I think it's wonderful" and "I think it is a great time" are positive, however it seems to have an obsession with talking about a "great time." This is probably because the regularization is not high enough, so maybe increasing α a bit would help, but overall the reward is solid and the KL-divergence stayed low, so it's a good balance.

(c) Based on your experience with REINFORCE with KL-regularization, identify one of its main limitations and propose a potential solution. There's no single right answer here; any justifiable answer will result in full credit.

One issue with REINFORCE with KL-regularization is that it can be difficult to tune the α hyperparameter. If α is too high, the model will be overly regularized and will not learn anything at all. If α is too low, the model will not be regularized at all, and will instead just try to maximize the reward by generating sequences that are likely to receive high rewards.

One potential solution to this issue is to use **annealing**. Basically, we start with a high value of α (high regularization) and then gradually decrease it over time. This can help to ensure that the model is regularized enough to prevent it from generating incoherent sequences, but is not overly regularized so that it cannot learn anything at all.

(d) Over which tokens do we compute the REINFORCE loss, and why?

We compute the REINFORCE loss over the tokens in the sequence themselves. This is because it allows us to take the gradient of the reward with respect to the model's parameters to favor sequences of tokens that lead to higher rewards. It's just basic gradient descent, but with the reward as the objective function. (Actually I guess in this case it's gradient ascent, since we're trying to maximize the reward.)




Figure 6. DPO training loss over time.

Question 4: DPO Implementation

(a) Run your training function for 2 epochs using the batch size 16. Plot the training losses over time and report the model's average loss on the test set for the final model. Note that an efficient implementation should run in about 10 minutes.

The DPO test loss was 0.5968.

(b) Run the evaluation function using the same data and code from Section 1 using the DPO-trained model. How does the reward compare to the REINFORCE model? Why do you think the results look like this?

There are many reasons this could happen, from hyperparameter sensitivity to the fact that REINFORCE can do reward hacking to get a really high reward. It's also possible that the DPO model is just not learning as well as the REINFORCE model, or that the DPO model is overfitting to the training data.

(c) Provide 5 example generations from the DPO-trained model. Are they more positive than the original generations? What else do you notice about them?

Five example generations are:

1. "Young Elijah Wood and Joseph Mazzello are outstanding musicians and performers for the United States, the United Kingdom, the Netherlands, Turkey, and France."
2. "This is one of the best sequels around and a real star, just because they are better than many others.\n\nYou may not know"
3. "That's what me and my friends kept asking each other the day about. At one point, I asked if she'd take some of it away from"
4. "'Mame' is a disgrace to many things he did.\nP.S. I don't actually like it too much...\nM"
5. "The Tender Hook, or, Who Killed The Tender Hook, or, Who Killed The Tender Hook, or, Who Killed the Tender"

These generations are not more positive than the original generations, and are in fact basically the same as the original generations. This is likely because the DPO model is just overfitting to the training data, or is really sensitive to the hyperparameters (a likely outcome as I haven't tuned them). Compared to REINFORCE, which had really different generations (but reward-hacked generations), the DPO model seems to be generating more natural language, even though the sentence doesn't make much sense. But the sentiment is not really positive.

(d) Try 3 different values for β (note it's typically set to be in the range of 0.1 to 0.5). Plot the training loss plots and report the final test losses for each of these setups. What's the best setup that you found? How do different values of β impact the training differently?

I tried β values of 0.1, 0.25, and 0.4.

The final test loss for $\beta = 0.1$ was 0.6521 (best setup).

The final test loss for $\beta = 0.25$ was 0.6207.

The final test loss for $\beta = 0.4$ was 0.6044.

It seems that as β increased, so did the test loss. The smaller values of β seem to perform better.

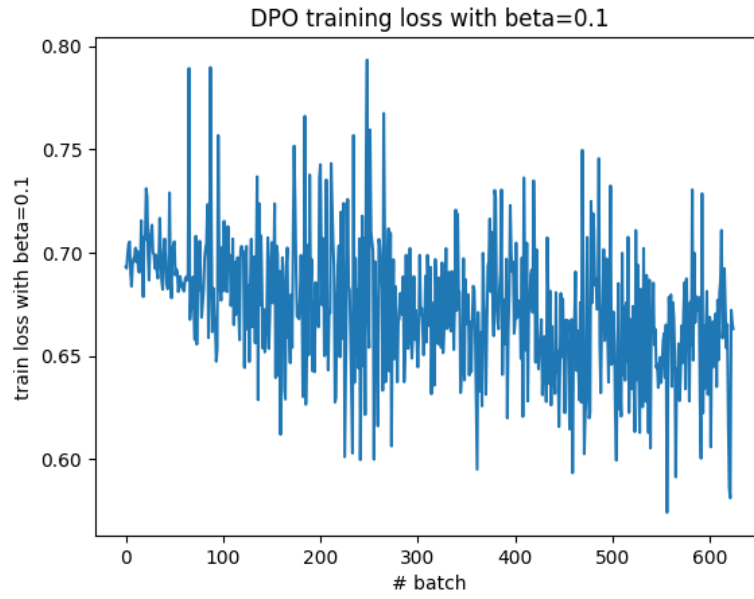


Figure 7. DPO training loss over time with $\beta = 0.1$.




Figure 8. DPO training loss over time with $\beta = 0.25$.




Figure 9. DPO training loss over time with $\beta = 0.4$.

(e) What's the purpose of using σ , the logistic function, in \mathcal{L}_{DPO} ?

σ is used to keep the output of the log probability ratios in the range $[0, 1]$. It also sort of smooths the output, so that the model can calculate gradients more easily.

(f) What are the main differences and advantages of DPO compared to other online RL methods like PPO?

DPO directly optimizes the policy, by increasing the likelihood of preferred actions and decreasing the likelihood of non-preferred actions. This is different from PPO, which uses a clipped objective function to prevent the policy from changing too much at each iteration.

Because PPO has a clipping mechanism, beneficial updates can be discarded if they fall outside the clipping range.

The logistic function in DPO provides smoother gradient updates as well (see above answer).

Question 5: GPT-4 Capability Forecast

Large language models are surprisingly good at some things, and surprisingly bad at others. Test and build your intuition for GPT-4's capabilities in this online quiz: <https://nicholas.carlini.com/writing/llm-forecast/question/Capital-of-Paris>

(a) Report your accuracy and log loss.

- **Accuracy:** 64.29%
- **Average Log Loss:** 1.079

This abysmal performance on my part suggests that I have a soft side for GPT-4, as I was wildly overconfident in my guesses, and thus my confidence levels were not well-calibrated.

(b) Was there anything you expected GPT-4 to be able to do that it couldn't do? Or vice versa?

I expected that GPT-4 would be able to answer certain questions like Password or Tic-Tac-Toe, but turns out these sorts of tasks are too complex or require logical reasoning that can't be handled.

On the other hand, I was surprised that it could do math really well (like take an integral) and also the JavaScript code was really cool! Making an entire flag was very impressive and something that would take me a really long time.

(c) What patterns do you notice in the failure/success modes? Can you guess the reasons behind these failure/success modes?

I noticed that GPT-4 was really good at generating text that was similar to the training data, but was bad at other things. For example, it was really easy to emulate the style of a certain author, but it was really hard to answer questions that required logical reasoning.

I guess the reasons for this are because GPT-4 (and language models) can mimic existing text really well, but actually forming a new idea (through logical reasoning) is not as easy, as there's no textual basis for it to draw from.

(d) Do you trust GPT-4 to help you with a task more or less after taking the quiz?

I trust GPT-4 less after taking the quiz. I was surprised at how bad it was at some things, and I think I would be more cautious about using it for tasks that require logical reasoning or understanding of the world.

(e) Using your wildest imagination, what capabilities do you hope the future GPT-x will have? Why are they important to have? What are some prerequisites if these capabilities were to be achieved in the future GPT-x? Be creative, and there's no single correct answer!

I hope that future GPT-x models will be able to understand the world and reason about it. This is important because it would allow us to use GPT-x for a wide range of tasks, from answering

questions to solving problems to generating new ideas.

I feel like a prerequisite for this would be safety. We need to make sure that if such an advanced language model would exist, it would not be used for malicious purposes. We also need to make sure that it is not biased, and that it is able to understand and reason about the world in a way that is consistent with human values.